

Computer Science for the Physical Sciences

Week 9

Craig Rasmussen (Research Support Services, University of Oregon)

Computer Science Minor: *This Week*

- Required courses (24 credits)
 - Introduction to Computer Science I-II-III
 - Elements of Discrete Mathematics I-II
 - Introduction to Data Structures
- Upper-division courses (8 credits)
 - Computer Architecture
 - Introduction to Algorithms
 - C/C++ and Unix nm
 - Operating Systems
 - Automata Theory
 - Software Methodology I-II Debugging
MPI Parallelism
 - Databases
 - Computational Science Numerical Differentiation
 - Bioinformatics
 - Data Mining
 - Introduction to Artificial Intelligence
 - Machine Learning

China: “Fastest” Computer in the World

34 Petaflops



ACISS System - Compute Resources

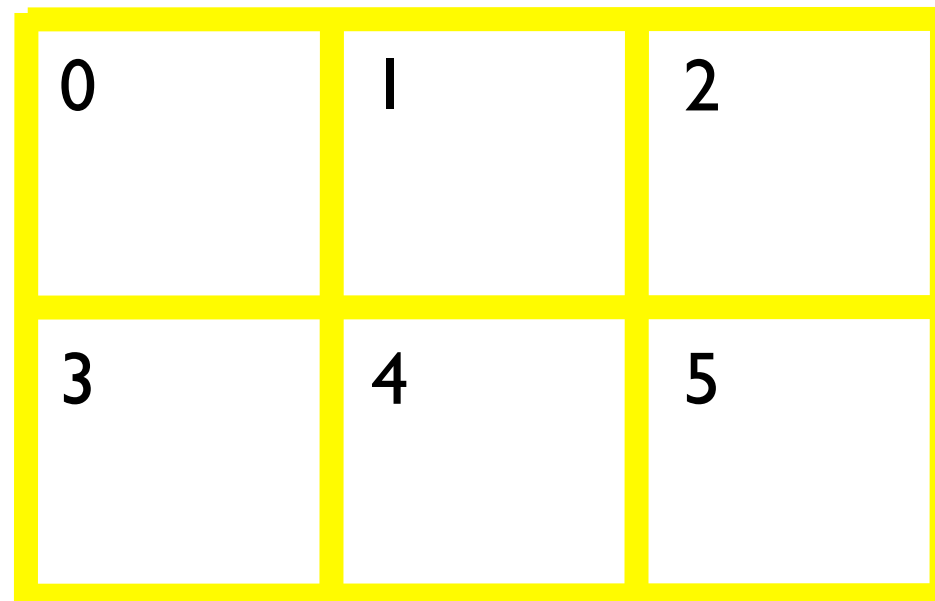
- Basic node (17 total Teraflops)
 - 128 ProLiant SL390 G7
 - Two Intel X5650 2.66 GHz 6-core CPUs per node (1,536 total cores)
 - 72 GB DDR3 RAM per basic node



Communication is paramount!

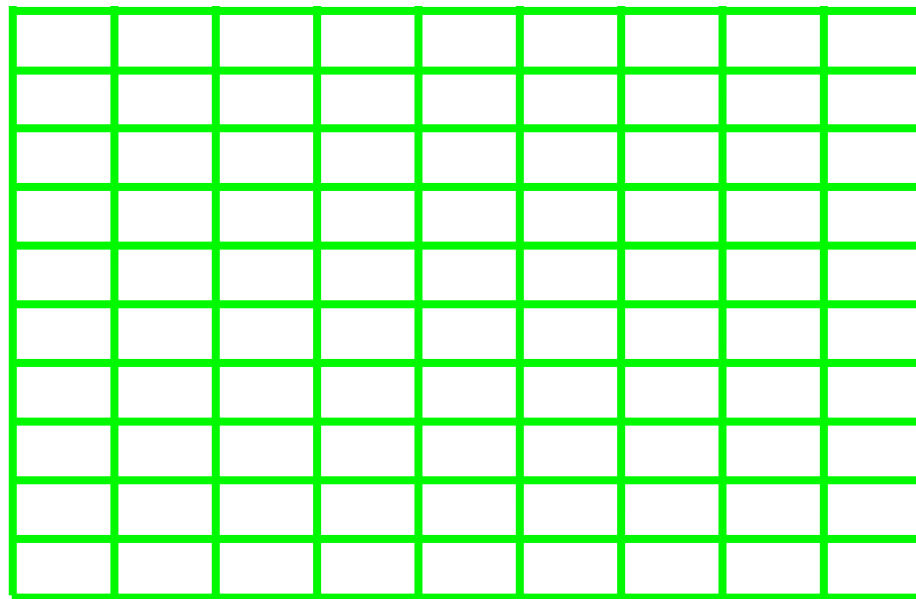
Processor Topology

- A set of P processes can be arranged in a virtual grid of $M_p * N_p = P$ processes
- 2 rows by 3 columns = 6 processes



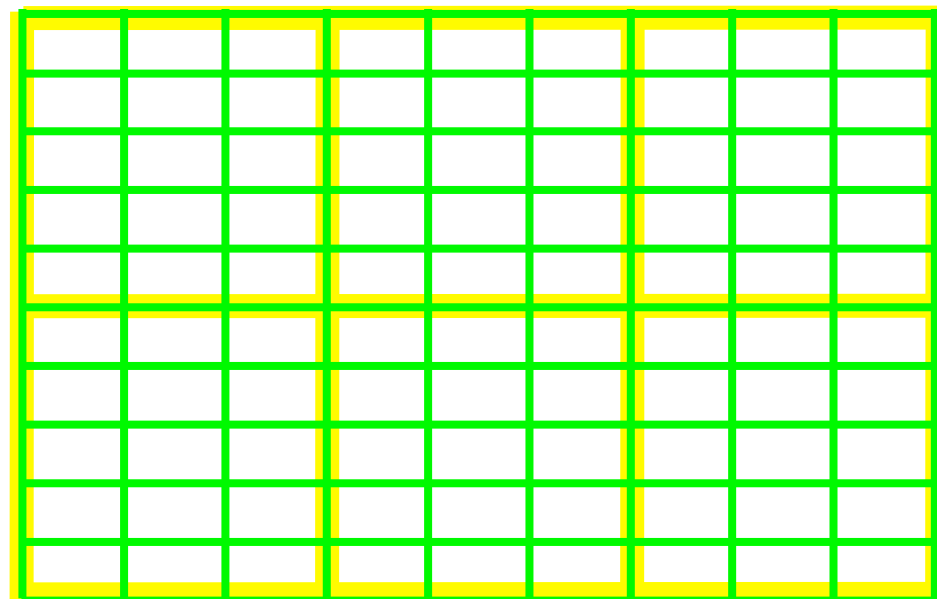
Data Grid Topology

- Cells of data can be arranged in a grid of $M \times N$
= L cells
- 10 rows by 9 columns = 90 cells



Data Parallelism (SPMD)

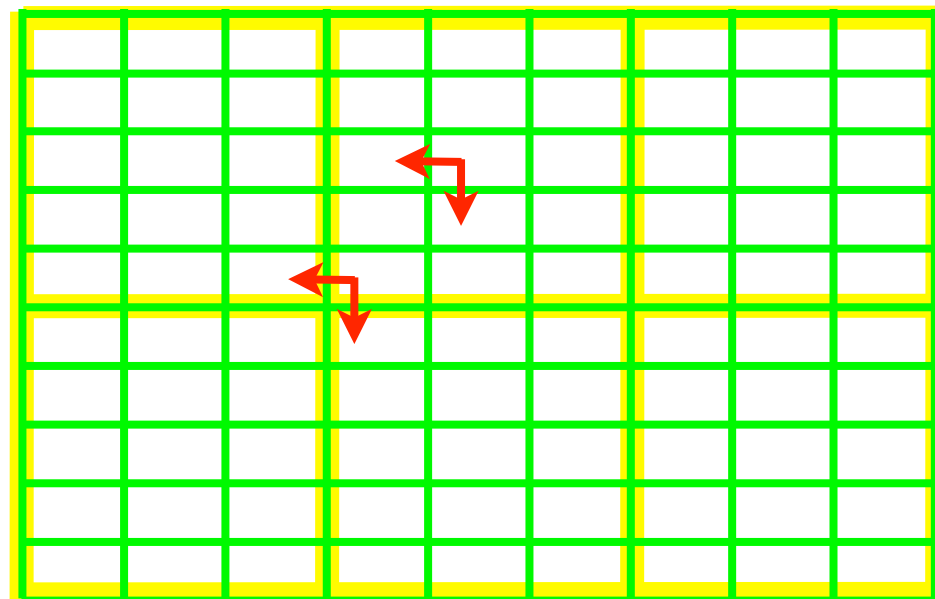
- Distribute a group of cells to a processor
- Scales well (up to 100 K + processors)
 - weak scaling: as cell count increases, use more processors
- 5 rows by 3 columns = 15 cells per processor



Data Parallelism: MPI

Communication

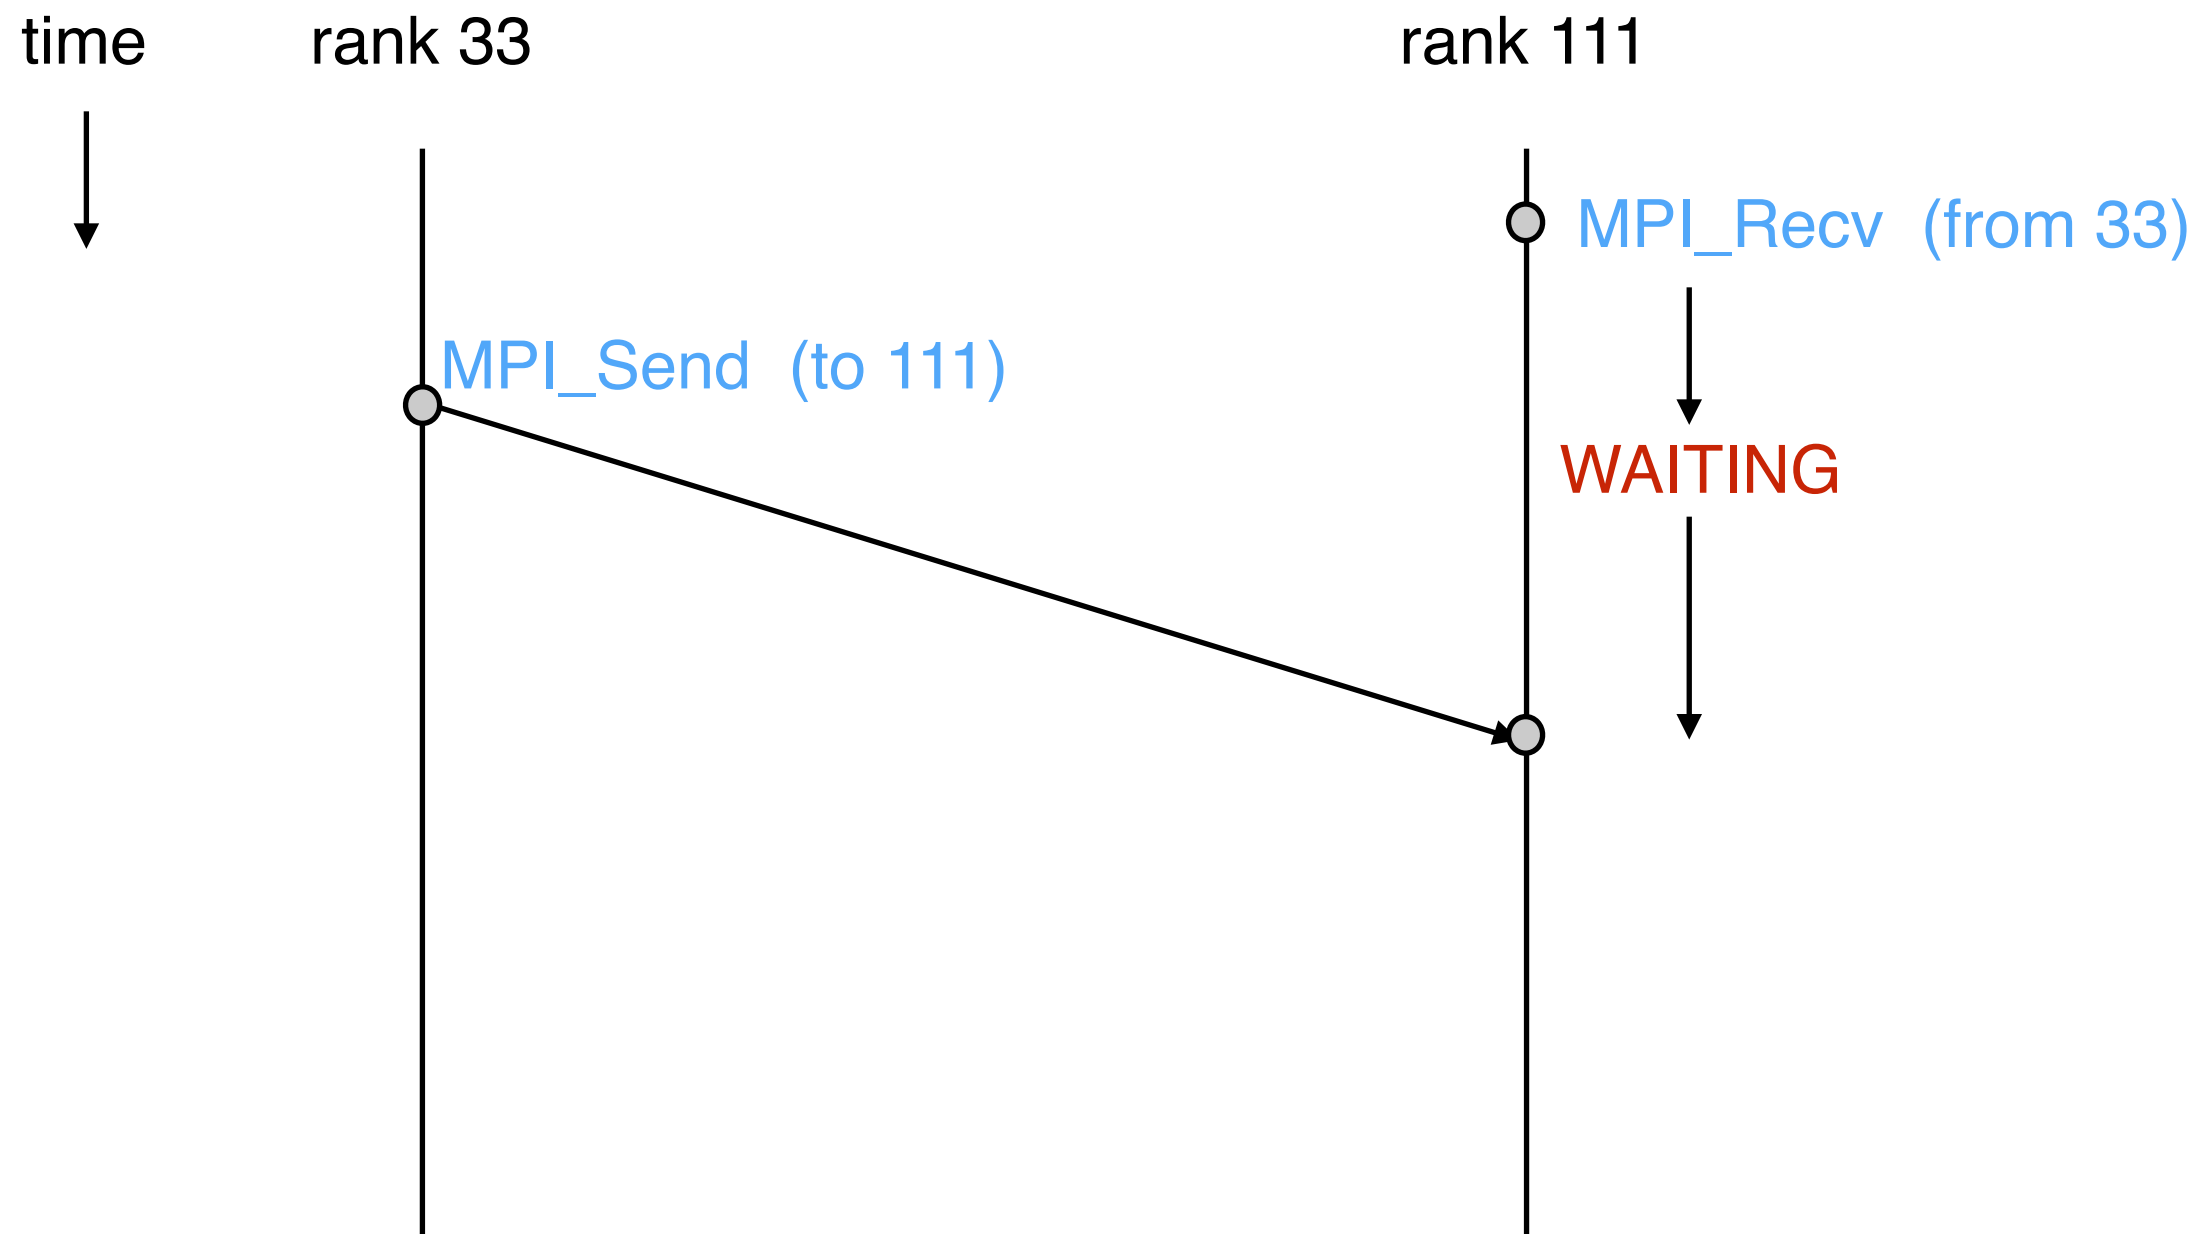
- Communication is usually between neighboring processors
 - communication costs scale as processor surface
 - computation done on volume (# of cells)
- Global communication often scales $O(\log(P))$
 - calculate an average over all cells (e.g., average temperature)



Parallelism with MPI: *An overview*

- MPI is remarkably easy
- Six function MPI
 - MPI_Init # initialize the MPI library
 - MPI_Comm_size # how many processes in the communicator
 - MPI_Comm_rank # what is my process # in the communicator
 - MPI_Send # send a buffer to a process (blocking)
 - MPI_Recv # receive from a process (blocking)
 - MPI_Finalize # finished using MPI
- MPI_COMM_WORLD is the default communicator that includes all processes
 - subsets of MPI_COMM_WORLD cab be taken

Parallelism with MPI: *Matching send and recv*



Parallelism with MPI: *Extra functions*

- More MPI functions
 - MPI_Barrier # **block** until all processes reach this point
 - MPI_Reduce # **sum reduction** (eg) of all elements in an array
 - MPI_Broadcast # **broadcast** a buffer to everyone
 - MPI_Scatter # **scatter sections** of a buffer to everyone
 - MPI_Gather # **gather sections** of a buffer from everyone

Shell Command: *nm*

- nm - display name list (symbol table)
- Steps to build an executable from a source file (*.f90)
 - *compile program* -> \$(FC) -c -I include_path hello.f90
 - *link program* -> \$(FC) -o hello hello.o -L library_path -lsome_library
- hello.o is an object file and contains symbols (e.g., functions) and code
- hello is an executable file (created by linker from *.o and libraries)
- What happens if a symbol (function code) can't be found by linker?
 - linker can't create an executable if all dependencies aren't satisfied
 - *use nm to track down missing symbols*

Debugging

- **ADVICE:** learn to use the debugger
- I make lots of errors in coding so I live in the debugger
- But sometimes with parallel programming printing is still a useful option for debugging